

Prasad Chandrashekhar Kanade

+1 6174125390 | kanade.pra@northeastern.edu | [linkedin/prasad-kanade](#) | [github/prasad0411](#) | [portfolio-website](#) | Boston, MA

EDUCATION

Master of Science in Computer Science

Northeastern University, Boston, USA

September 2025 - May 2027

3.73/4

Courses: Database Management, Algorithms, Machine Learning, Distributed Systems, Cloud Computing

Bachelor of Technology in Computer Science and Engineering

Maharashtra Institute of Technology - World Peace University, Pune, India

October 2020 - November 2023

9.53/10

Courses: Data Structures, Probability & Statistics, Software Engineering, Computer Networks, Artificial Intelligence

TECHNICAL SKILLS

- **Languages:** Python, SQL, Shell/Bash, R, Java
- **Pipelines & Orchestration:** Apache Spark (PySpark), Kafka, Airflow, dbt, ETL/ELT, incremental processing, streaming ingestion, DAG scheduling, backfills
- **Warehousing & Modeling:** Snowflake, BigQuery, Databricks, Delta Lake, medallion architecture, dimensional modeling, staging and mart layers, PostgreSQL, Oracle, MySQL, MongoDB, Redis, ChromaDB
- **Data Quality:** Schema validation, contract testing, deduplication, entity resolution, anomaly detection, statistical process control, drift monitoring, lineage tracking
- **Cloud & Analytics:** AWS (EC2, ECR, ECS), Azure Data Factory, ADLS, Docker, Terraform, GitHub Actions, Linux, CI/CD, Pandas, NumPy, Tableau, Power BI, Streamlit

WORK EXPERIENCE

Amdocs

December 2023 - May 2025

Data Engineer

Pune, India

- Engineered event-driven pipelines processing 730M+ events annually (2M+ daily across 10M+ accounts), with Kafka handling ingestion and Airflow DAGs orchestrating validation, transformation, and analytical workloads with backfills.
- Built 22 dbt models across staging and mart layers over subscriber, account, and event data, with 38 tests on uniqueness, nullability, referential integrity, accepted values, and schema consistency - delete-and-insert cut runtime 65%.
- Integrated Snowflake for a ~750 GB dataset alongside Oracle, moving analytical workloads to Delta Lake Bronze/Silver/Gold on PySpark/Databricks - incremental processing avoided table rebuilds, cutting transformation 50%.
- Enforced schema contracts and statistical anomaly detection across 80% of the pipeline, quarantining malformed records and detecting distribution shift before downstream consumers - cutting production data incidents 30%.

ISKCON

July 2022 - December 2022

Data Engineer Intern

Pune, India

- Automated collection, validation, and reconciliation across 50,000+ participant records spanning five departments through scheduled Python and SQL workflows, replacing manual entry and producing analysis-ready datasets.
- Built reporting workflows over survey and operational data collected through Qualtrics and Google Forms across 3,000+ respondents, standardizing inconsistent inputs into consistent schemas for downstream analysis.

Simba Developers

March 2021 - October 2021

Backend Engineer Intern

Pune, India

- Optimized data-intensive application performance 44% (3.2s to 1.8s) by profiling query hot paths, caching repeated lookups in Redis, and deferring non-critical computation off the request path.
- Built PostgreSQL-backed services processing real-time payment transaction data through Node.js and Express, with retry policies and resource isolation that reduced cascade failures 65% when downstream dependencies degraded.

PROJECTS

Job Market Intelligence Platform [GitHub](#)

February 2026 - Present

- Built a PySpark medallion lakehouse over 25 source schemas landing Bronze, Silver, and Gold Parquet layers, with dbt staging and mart models orchestrated through Airflow DAGs and gated by 35 data tests in CI.
- Partitioned by source and ISO week, cutting single-week query scans 86% (752 KB of 5.5 MB) with file-level pruning confirmed in the Spark physical plan, and incremental materialization rebuilding only the affected window.

Thyroid Disease Classification - Clinical Dashboard [GitHub](#) | [Live Demo](#) | [Paper](#)

September 2025 - May 2026

- Built a clinical analytics platform over a 150,000-record dataset spanning feature engineering, batch prediction, drift detection on incoming distributions, and SHAP attribution, served through Streamlit and FastAPI.
- Reached 85.1% accuracy over a 73% baseline with leakage-audited feature selection and per-class recall published.