

Prasad Chandrashekhar Kanade

+1 6174125390 | kanade.pra@northeastern.edu | [linkedin/prasad-kanade](https://www.linkedin.com/in/prasad-kanade) | [github/prasad0411](https://github.com/prasad0411) | [portfolio-website](#) | Boston, MA

EDUCATION

Master of Science in Computer Science

Northeastern University, Boston, USA

September 2025 - May 2027

3.73/4

Courses: Machine Learning, Distributed Systems, Database Management, Algorithms, Cloud Computing

Bachelor of Technology in Computer Science and Engineering

Maharashtra Institute of Technology - World Peace University, Pune, India

October 2020 - November 2023

9.53/10

Courses: Artificial Intelligence, Probability and Statistics, Data Structures, Software Engineering, Computer Networks

TECHNICAL SKILLS

- **Languages:** Python, SQL, Java, C++, R, Shell/Bash
- **Machine Learning:** PyTorch, scikit-learn, XGBoost, Random Forest, ensembles, Optuna, SMOTE, RFE, SHAP, probability calibration, imbalanced learning, cross-validation design, leakage auditing
- **LLM & Generative AI:** RAG, LangGraph, LoRA/PEFT fine-tuning, embeddings, semantic search, retrieval evaluation, prompt engineering, agent workflows
- **MLOps & Data:** MLflow, experiment tracking, model versioning, FastAPI, batch and online inference, drift detection, model cards, Docker, GitHub Actions, Kafka, Pandas, NumPy, feature pipelines, anomaly detection
- **Storage:** ChromaDB, FAISS, PostgreSQL, Oracle, Redis

WORK EXPERIENCE

Amdocs

December 2023 - May 2025

Machine Learning Systems Engineer

Pune, India

- Engineered fault-tolerant feature delivery across 8 Kafka services streaming events for 10M+ accounts, with retries, circuit breakers, and graceful degradation guaranteeing ML consumers fresh features at 99.5% uptime.
- Developed statistical anomaly detection in Python with threshold and distribution-based monitoring, flagging feature drift and schema changes in real time and cutting data incidents 30% before corrupted input reached downstream models.
- Reduced feature-retrieval latency 60% by profiling Oracle execution plans, introducing composite indexes, and caching high-frequency lookups in Redis - removing 15+ hours/week of batch delay from the feature-preparation cycle.
- Delivered validated feature datasets consumed by predictive and reporting workflows, enforcing schema contracts at ingest across 80% of the pipeline so downstream models read consistent inputs rather than source-shaped records.

ISKCON

July 2022 - December 2022

Data & Infrastructure Engineer Intern

Pune, India

- Automated data collection, validation, and reconciliation across 50K+ participant records, replacing manual entry across five departments and producing analysis-ready datasets for downstream reporting.
- Instrumented Unix infrastructure with Prometheus, Grafana, and Alertmanager, threshold-alerting on time-series metrics for CPU, memory, and disk - cutting mean-time-to-detection 4+ hours/week and unplanned downtime 25%.

Simba Developers

March 2021 - October 2021

Applied Machine Learning Engineer Intern

Pune, India

- Built ML-driven irregularity detection over transaction data to flag anomalous payment behavior, owning data preparation, feature selection, model evaluation, and inference integration into the Node.js/Express request path.
- Served predictions through a PostgreSQL and Redis layer feeding WebSocket notifications, isolating inference behind retry policies and bulkheads so model failures could not cascade into the payment flow - cutting cascade failures 65%.

PROJECTS

Thyroid Disease Classification - Clinical Decision Support [GitHub](#) | [Live Demo](#) | [Paper](#) September 2025 - May 2026

- Built a clinical decision support system with SHAP explainability and a RAG pipeline benchmarked across NumPy, FAISS, and ChromaDB backends, plus LoRA/PEFT fine-tuning and LangGraph agent orchestration tracked in MLflow.
- Trained XGBoost on 150,000 records to 85.1% accuracy and macro F1 0.727 ± 0.004 , with SMOTE inside CV folds and leakage-audited feature selection; served via FastAPI at 0.54 ms, 11,700 predictions/second.

Job Market Intelligence Platform [GitHub](#)

February 2026 - Present

- Engineered a data platform ingesting 15,000+ records weekly across 25 sources, combining text classification, semantic similarity, and automated quality correction into a pipeline that runs unattended.
- Designed a self-correcting entity resolution layer using TF-IDF similarity, fuzzy matching, and a confidence cache that improves with every correction, plus monitoring that isolates degrading sources before contamination.