

Prasad Chandrashekhar Kanade

+1 6174125390 | kanade.pra@northeastern.edu | [linkedin/prasad-kanade](#) | [github/prasad0411](#) | [portfolio-website](#) | Boston, MA

EDUCATION

Master of Science in Computer Science

Northeastern University, Boston, USA

September 2025 - May 2027

3.73/4

Courses: Distributed Systems, Database Management, Algorithms, Cloud Computing, Machine Learning

Bachelor of Technology in Computer Science and Engineering

Maharashtra Institute of Technology - World Peace University, Pune, India

October 2020 - November 2023

9.53/10

Courses: Data Structures, Operating Systems, Software Engineering, Computer Networks, Artificial Intelligence

TECHNICAL SKILLS

- **Languages:** Java, Python, TypeScript, SQL, C++, Shell/Bash
- **Backend & APIs:** Spring Boot, REST, gRPC, Kafka, JPA/Hibernate, Spring Security, JWT, OAuth2, OpenAPI/Swagger, Node.js, Express.js, FastAPI, Microservices, LangGraph, RAG, vector databases
- **Cloud & DevOps:** AWS (EC2, ECR, ECS), Docker, Terraform, GitHub Actions, Linux, Prometheus, Grafana
- **Databases:** Oracle, PostgreSQL, MySQL, MongoDB, Redis, SQLite, ChromaDB, FAISS
- **Testing & Quality:** JUnit, Mockito, pytest, contract testing, static analysis, automated quality gates

WORK EXPERIENCE

Amdocs

December 2023 - May 2025

Software Engineer

Pune, India

- Owned 8 Spring Boot microservices serving 10M+ subscribers and consuming 2M+ Kafka events daily, partitioning topics across consumer groups, containing failures behind circuit breakers and bulkhead isolation to hold 99.5% uptime.
- Cut Oracle query latency 60% by profiling AWR execution plans to isolate access paths, introducing composite indexes, and adding Redis caching for high-frequency subscriber lookups - removing 15+ hours/week batch delay.
- Established automated testing across all 8 services with JUnit and Mockito spanning unit, integration, and API contract layers, reaching 80% coverage and enabling regression gating on every release.
- Standardized API contracts in OpenAPI across REST and gRPC services and introduced team-wide review conventions, cutting integration defects 40% and shortening the release cycle from two weeks to four days.
- Built data validation and anomaly detection in Python and Shell across telecom pipelines, catching schema drift and malformed records before they reached downstream consumers and cutting production data incidents 30%.

ISKCON

July 2022 - December 2022

Software Engineer Intern

Pune, India

- Instrumented Unix infrastructure with Prometheus, Grafana, Node Exporter, and Alertmanager covering CPU, memory, disk, and service health, cutting mean-time-to-detection 4+ hours/week and unplanned downtime 25%.
- Containerized and deployed workloads to AWS EC2 via Docker, ECR, and ECS on Terraform-managed infrastructure, building CI/CD pipelines that reduced deployment time 78% and eliminated manual release steps.

Simba Developers

March 2021 - October 2021

Software Developer Intern

Pune, India

- Built fault-tolerant Node.js/Express services backed by PostgreSQL and Redis, using WebSockets for real-time payment notifications and isolating dependencies behind retry policies and resource partitioning to reduce cascade failures 65%.
- Improved data-intensive application response time 44% (3.2s to 1.8s) by profiling hot paths to isolate bottlenecks, caching repeated lookups in Redis, and deferring non-critical render work.

PROJECTS

Job Market Intelligence Platform [GitHub](#)

February 2026 - Present

- Designed a platform integrating 10 job-board APIs behind per-provider isolation and exponential backoff, quarantining failing upstreams rather than stalling the run - ingesting 15,000+ postings weekly across 25 sources.
- Made concurrent persistence crash-safe across five processes sharing one datastore through file locking, atomic replace, and write-ahead journaling rather than a central lock service, so interrupted runs resume without manual recovery.

Thyroid Disease Classification - Clinical Decision Support [GitHub](#) | [Live Demo](#) | [Paper](#)

September 2025 - May 2026

- Built a clinical decision support system that classifies thyroid conditions from lab panels, explains each prediction with SHAP attributions, and answers clinician follow-ups through a RAG pipeline over an indexed medical corpus.
- Trained XGBoost on 150,000 records reaching 85.1% accuracy against a 73% baseline, with leakage checks and cross-validated feature selection; served predictions through FastAPI in under 1 ms, scaling to 11,700/second batched.